

Predicting a Binary Classifier

The purpose of this project is to create a machine learning model to predict whether entries in the dataset belong to a specific class (' y '). This dataset was provided by a client as skills assessment, and the objectives and guidelines for the project were outlined by the client in advance. The steps required to complete the project will be:

- Data Cleansing
- Feature Engineering and Selection
- Building two different models on training dataset
- Model tuning
- Generating predictions for test dataset
- Comparing modeling approaches in a writeup
- Preparing all files for submission

Initial Analysis and Data Cleansing

To begin, I will import the 'training' dataset into a pandas dataframe and perform initial analysis to determine what cleansing steps need to be taken.

```
In [2]: # Importing necessary data manipulation/visualization libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

# Expanding number of rows and columns allowed in view
pd.set_option('max_columns', 180)
pd.set_option('max_rows', 200000)
pd.set_option('max_colwidth', 5000)
```

```
In [3]: # Reading training and test data into dataframes
train = pd.read_csv('traindata.csv')
test = pd.read_csv('testdata.csv')
```

```
In [4]: # Exploring shape and head of test data
print(test.shape)
test.head()
```

```
(10000, 100)
```

Out[4]:

	x0	x1	x2	x3	x4	x5	x6	x7
0	0.519093	-4.606038	13.707586	-17.990903	12.873394	14.910935	2.915341	-10.110081
1	-12.357004	13.874141	14.052924	34.129247	34.511107	34.583336	-0.482540	-6.583407
2	1.834922	2.665252	-44.873210	21.941920	10.102981	5.962249	-5.733909	-4.061670
3	20.972483	11.548506	-40.924625	-35.296796	-35.253101	-14.601890	5.045075	10.841771
4	-9.916044	5.509811	31.749288	-0.803916	-4.005098	20.912490	0.419346	-2.949516

```
In [5]: # Exploring shape and head of train data
print(train.shape)
train.head()
```

```
(40000, 101)
```

Out[5]:

	x0	x1	x2	x3	x4	x5	x6	x7
0	0.963686	6.627185	-45.224008	9.477531	-3.216532	13.216874	9.754747	5.245851
1	-1.770062	-23.610459	-0.964003	-31.981497	-10.294599	-10.240251	-1.518888	-1.675208
2	9.962401	-8.349849	23.248891	-24.196879	8.937480	10.965000	-7.490596	-3.025094
3	-5.780709	-25.261584	1.383115	-11.786929	7.993078	-11.245752	-2.607351	-3.513896
4	1.211541	1.119963	7.512938	21.987312	-5.155392	10.339416	3.045180	-0.619230

```
In [6]: # Analyzing quick statistics on distribution of each feature
train.describe()
```

Out[6]:

	x0	x1	x2	x3	x4	x5
count	39988.000000	39990.000000	39993.000000	39987.000000	39993.000000	39992.000000
mean	2.020255	-3.924559	1.006619	-1.378330	0.070199	-0.715213
std	9.590599	18.768656	21.062970	29.397779	20.243287	18.268807
min	-36.842503	-79.156374	-89.728356	-126.652341	-76.412886	-73.743342
25%	-4.461433	-16.591552	-13.230956	-21.297149	-13.580632	-13.092873
50%	2.022412	-4.061703	1.184946	-1.224625	0.091600	-0.657601
75%	8.389979	8.529110	15.221205	18.530623	13.722427	11.610239
max	44.478690	77.682652	84.625640	117.004453	85.934044	74.465465

```
In [7]: # Looking at types of each feature  
train.info(verbose=True)
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 40000 entries, 0 to 39999  
Data columns (total 101 columns):  
x0      float64  
x1      float64  
x2      float64  
x3      float64  
x4      float64  
x5      float64  
x6      float64  
x7      float64  
x8      float64  
x9      float64  
x10     float64  
x11     float64  
x12     float64  
x13     float64  
x14     float64  
x15     float64  
x16     float64  
x17     float64  
x18     float64  
x19     float64  
x20     float64  
x21     float64  
x22     float64  
x23     float64  
x24     float64  
x25     float64  
x26     float64  
x27     float64  
x28     float64  
x29     float64  
x30     float64  
x31     float64  
x32     float64  
x33     float64  
x34     object  
x35     object  
x36     float64  
x37     float64  
x38     float64  
x39     float64  
x40     float64  
x41     object  
x42     float64  
x43     float64  
x44     float64  
x45     object  
x46     float64  
x47     float64  
x48     float64  
x49     float64  
x50     float64  
x51     float64  
x52     float64  
x53     float64
```

```
x54    float64
x55    float64
x56    float64
x57    float64
x58    float64
x59    float64
x60    float64
x61    float64
x62    float64
x63    float64
x64    float64
x65    float64
x66    float64
x67    float64
x68    object
x69    float64
x70    float64
x71    float64
x72    float64
x73    float64
x74    float64
x75    float64
x76    float64
x77    float64
x78    float64
x79    float64
x80    float64
x81    float64
x82    float64
x83    float64
x84    float64
x85    float64
x86    float64
x87    float64
x88    float64
x89    float64
x90    float64
x91    float64
x92    float64
x93    object
x94    float64
x95    float64
x96    float64
x97    float64
x98    float64
x99    float64
y      int64
dtypes: float64(94), int64(1), object(6)
memory usage: 30.8+ MB
```

```
In [8]: # Checking null counts of each feature  
nullcounts = train.isnull().sum()  
nullcounts.sort_values(ascending=False, inplace=True)  
nullcounts
```

```
Out[8]: x55      16
        x13      15
        x18      14
        x96      14
        x3       13
        x99      13
        x51      13
        x62      13
        x73      13
        x21      12
        x85      12
        x0       12
        x60      12
        x56      11
        x24      11
        x17      11
        x33      11
        x69      11
        x42      11
        x12      11
        x77      11
        x65      10
        x7       10
        x59      10
        x39      10
        x35      10
        x26      10
        x1       10
        x66      10
        x89      10
        x75      10
        x28       9
        x11       9
        x23       9
        x46       9
        x94       9
        x6        9
        x87       9
        x58       9
        x97       9
        x31       9
        x68       9
        x72       9
        x74       9
        x19       8
        x15       8
        x78       8
        x5        8
        x48       8
        x82       8
        x16       8
        x34       8
        x63       8
        x40       8
        x76       8
        x57       8
        x25       7
```

```
x90      7
x80      7
x4        7
x45      7
x79      7
x10      7
x27      7
x95      7
x52      7
x93      7
x61      7
x86      7
x2        7
x67      7
x9        7
x22      6
x70      6
x36      6
x38      6
x71      6
x8        6
x92      6
x54      6
x14      5
x50      5
x53      5
x88      5
x37      5
x64      5
x20      4
x98      4
x83      4
x49      4
x41      4
x47      3
x29      3
x44      3
x81      3
x84      3
x32      3
x91      2
x30      2
x43      1
y         0
dtype: int64
```

```
In [9]: # Printing total of all null values
print(nullcounts.sum())
nullpct = nullcounts.sum()/len(train)
print(nullpct)
```

```
806
0.02015
```


Initial Observations

After a quick exploration of the training dataset, a few key observations stand out:

- The data is comprised of 40,000 entries, with 100 features.
- A small portion of the features are string objects, with the rest stored as numeric (float) features
- The features appear to be of different ranges and sizes, but based on the min and max values and standard deviations there do not appear to be any huge outliers.
- Each feature has some null values, with only 16 in any one column and 806 null values overall. This is a small percentage of the overall dataset (2.01% of total data).
- The 'y' , or target, column is binary. Each entry is listed as a 0 or 1 in this column, and there are no null values.
- The target column is unbalanced. Approximately 20% of the data is categorized as a '1' and 80% as a '0' . I will need to account for this imbalance in the machine learning models.

Cleansing Steps Needed

To prepare the data for model training and testing, I need to cleanse the data. This will include:

- Examining and standardizing the text-based features
- Converting text-based features to dummy variables
- Handling null values for each feature
- Standardizing all columns with a min-max scaler for efficient machine learning input

Note: Because these same processes will need to be performed on the testing dataset, the data cleansing will need to be incorporated into a function that can be applied to the test dataset at a later point.

Data Cleansing

Text-Based Features

```
In [10]: # Identifying text columns
text_vals = train.select_dtypes(include=['object'])
text_vals.head()
```

Out[10]:

	x34	x35	x41	x45	x68	x93
0	chrystler	thur	\$-865.28	0.02%	sept.	asia
1	volkswagon	thur	\$325.27	-0.01%	July	asia
2	bmw	thursday	\$743.91	0.0%	July	asia
3	nissan	thursday	\$538.48	0.01%	July	asia
4	volkswagon	wed	\$-433.65	0.0%	Jun	asia

```
In [11]: # Cleaning the values of x41 and x45, storing as float
text_vals['x41'] = text_vals['x41'].str.replace('$', '').astype('float')
text_vals['x45'] = text_vals['x45'].str.replace('%', '').astype('float')
```

/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>
This is separate from the ipykernel package so we can avoid doing imports until

```
In [12]: # Checking value counts of x34
text_vals['x34'].value_counts(sort=True)
```

```
Out[12]: volkswagon    12622
Toyota      10968
bmw         7262
Honda       5174
tesla       2247
chrystler   1191
nissan       326
ford        160
mercades    31
chevrolet   11
Name: x34, dtype: int64
```

These values seem to be fairly standardized, and are not significantly skewed. Several contain spelling errors which can be quickly cleaned with a mapping dictionary.

```
In [13]: # Creating a mapping dictionary for correct labeling
dict = {'volkswagon': 'volkswagen', 'Toyota': 'toyota',
        'bmw': 'bmw', 'Honda': 'honda', 'tesla': 'tesla',
        'chrystler': 'chrysler', 'nissan': 'nissan',
        'ford': 'ford', 'mercedes': 'mercedes',
        'chevrolet': 'chevrolet'}
# Replacing values using dictionary
text_vals['x34'] = text_vals['x34'].replace(dict)

/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:8: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy
```

```
In [14]: # Checking value counts of x35
text_vals['x35'].value_counts(sort=True)
```

```
Out[14]: wed          14820
thursday       13324
wednesday        5938
thur           4428
tuesday         884
friday          517
monday          53
fri             26
Name: x35, dtype: int64
```

Again, these are not significantly skewed, but several are misspelled and miscategorized. This can be quickly cleaned.

```
In [15]: # Creating mapping dictionary
dict = {'wed': 'wed', 'thursday': 'thu',
        'wednesday': 'wed', 'thur': 'thu',
        'tuesday': 'tue', 'friday': 'fri',
        'monday': 'mon', 'fri': 'fri'}
# Replacing values using dictionary
text_vals['x35'] = text_vals['x35'].replace(dict)

/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:7: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy
import sys
```

```
In [16]: # Checking value counts of x68
text_vals['x68'].value_counts(sort=True)
```

```
Out[16]: July          11114
Jun           9317
Aug           8170
May           4744
sept.         3504
Apr           1629
Oct           885
Mar           407
Nov           145
Feb           48
Dev           16
January       12
Name: x68, dtype: int64
```

These can again be cleaned with simple mapping. The one exception is the "Dev" value, which I am assuming is a misspelling of December.

```
In [17]: # Creating mapping dictionary
dict = {'July': 'jul', 'Jun': 'jun',
        'Aug': 'aug', 'May': 'may',
        'sept.': 'sep', 'Apr': 'apr',
        'Oct': 'oct', 'Mar': 'mar',
        'Nov': 'nov', 'Feb': 'feb',
        'Dev': 'dec', 'January': 'jan'}
# Replacing values using dictionary
text_vals['x68'] = text_vals['x68'].replace(dict)
```

```
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:9: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy
if __name__ == '__main__':
```

```
In [18]: # Checking value counts of x93
text_vals['x93'].value_counts(sort=True)
```

```
Out[18]: asia          35384
america          3167
euorpe           1442
Name: x93, dtype: int64
```

This feature seems to be heavily weighted in one category 'asia', representing 88% of the total entries. This kind of heavily skewed feature can be problematic for the machine learning model, so this feature should be dropped.

```
In [19]: # Drop x93 from text vals dataframe
text_vals = text_vals.drop(columns='x93')
```

```
In [20]: text_vals.head()
```

Out[20]:

	x34	x35	x41	x45	x68
0	chrysler	thu	-865.28	0.02	sep
1	volkswagen	thu	325.27	-0.01	jul
2	bmw	thu	743.91	0.00	jul
3	nissan	thu	538.48	0.01	jul
4	volkswagen	wed	-433.65	0.00	jun

Now that I have cleaned and standardized the text columns, I want to transform them into dummy categories to ensure the machine learning model is operating on strictly numeric data.

```
In [21]: # Selecting only text columns
textcols = ['x34', 'x35', 'x68']
# Get dummy prefix names
prefixes = list(text_vals[textcols].columns)
# Getting dummy features for each text column
text_dummies = pd.get_dummies(text_vals, prefix=prefixes)
```

```
In [22]: # Dropping dirty text columns from original dataset
dropcols = ['x34', 'x35', 'x41', 'x45', 'x68', 'x93']
train = train.drop(columns=dropcols)
```

```
In [23]: # Reintroducing our cleaned data into the original dataset
train = pd.concat([train, text_dummies], axis=1)
```

```
In [24]: train.info(verbose=True)
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 40000 entries, 0 to 39999  
Data columns (total 124 columns):  
x0                float64  
x1                float64  
x2                float64  
x3                float64  
x4                float64  
x5                float64  
x6                float64  
x7                float64  
x8                float64  
x9                float64  
x10               float64  
x11               float64  
x12               float64  
x13               float64  
x14               float64  
x15               float64  
x16               float64  
x17               float64  
x18               float64  
x19               float64  
x20               float64  
x21               float64  
x22               float64  
x23               float64  
x24               float64  
x25               float64  
x26               float64  
x27               float64  
x28               float64  
x29               float64  
x30               float64  
x31               float64  
x32               float64  
x33               float64  
x36               float64  
x37               float64  
x38               float64  
x39               float64  
x40               float64  
x42               float64  
x43               float64  
x44               float64  
x46               float64  
x47               float64  
x48               float64  
x49               float64  
x50               float64  
x51               float64  
x52               float64  
x53               float64  
x54               float64  
x55               float64  
x56               float64  
x57               float64
```

x58	float64
x59	float64
x60	float64
x61	float64
x62	float64
x63	float64
x64	float64
x65	float64
x66	float64
x67	float64
x69	float64
x70	float64
x71	float64
x72	float64
x73	float64
x74	float64
x75	float64
x76	float64
x77	float64
x78	float64
x79	float64
x80	float64
x81	float64
x82	float64
x83	float64
x84	float64
x85	float64
x86	float64
x87	float64
x88	float64
x89	float64
x90	float64
x91	float64
x92	float64
x94	float64
x95	float64
x96	float64
x97	float64
x98	float64
x99	float64
y	int64
x41	float64
x45	float64
x34_bmw	uint8
x34_chevrolet	uint8
x34_chrysler	uint8
x34_ford	uint8
x34_honda	uint8
x34_mercedes	uint8
x34_nissan	uint8
x34_tesla	uint8
x34_toyota	uint8
x34_volkswagen	uint8
x35_fri	uint8
x35_mon	uint8
x35_thu	uint8
x35_tue	uint8


```
x35_wed      uint8
x68_apr      uint8
x68_aug      uint8
x68_dec      uint8
x68_feb      uint8
x68_jan      uint8
x68_jul      uint8
x68_jun      uint8
x68_mar      uint8
x68_may      uint8
x68_nov      uint8
x68_oct      uint8
x68_sep      uint8
dtypes: float64(96), int64(1), uint8(27)
memory usage: 30.6 MB
```

We now have a dataframe including only numeric values.

Handling Missing Values

I will now correct any missing values. The amount of missing data is small in proportion to the overall dataset. I could easily replace null values with the mean of each feature. However, since I do not know the nature of each feature and what it represents, this could result in distorted, noisy information that skews the overall results of the model. Instead, I will simply drop any rows that contain null values.

```
In [25]: # Dropping rows with null values
train = train.dropna().copy()
```

```
In [26]: nullcounts = train.isnull().sum()  
nullcounts
```

```
Out[26]: x0      0
          x1      0
          x2      0
          x3      0
          x4      0
          x5      0
          x6      0
          x7      0
          x8      0
          x9      0
          x10     0
          x11     0
          x12     0
          x13     0
          x14     0
          x15     0
          x16     0
          x17     0
          x18     0
          x19     0
          x20     0
          x21     0
          x22     0
          x23     0
          x24     0
          x25     0
          x26     0
          x27     0
          x28     0
          x29     0
          x30     0
          x31     0
          x32     0
          x33     0
          x36     0
          x37     0
          x38     0
          x39     0
          x40     0
          x42     0
          x43     0
          x44     0
          x46     0
          x47     0
          x48     0
          x49     0
          x50     0
          x51     0
          x52     0
          x53     0
          x54     0
          x55     0
          x56     0
          x57     0
          x58     0
          x59     0
          x60     0
```

x61	0
x62	0
x63	0
x64	0
x65	0
x66	0
x67	0
x69	0
x70	0
x71	0
x72	0
x73	0
x74	0
x75	0
x76	0
x77	0
x78	0
x79	0
x80	0
x81	0
x82	0
x83	0
x84	0
x85	0
x86	0
x87	0
x88	0
x89	0
x90	0
x91	0
x92	0
x94	0
x95	0
x96	0
x97	0
x98	0
x99	0
y	0
x41	0
x45	0
x34_bmw	0
x34_chevrolet	0
x34_chrysler	0
x34_ford	0
x34_honda	0
x34_mercedes	0
x34_nissan	0
x34_tesla	0
x34_toyota	0
x34_volkswagen	0
x35_fri	0
x35_mon	0
x35_thu	0
x35_tue	0
x35_wed	0
x68_apr	0
x68_aug	0

```
x68_dec      0
x68_feb      0
x68_jan      0
x68_jul      0
x68_jun      0
x68_mar      0
x68_may      0
x68_nov      0
x68_oct      0
x68_sep      0
dtype: int64
```

```
In [27]: print(len(train))
         print(len(train[train['y']==1]))
```

```
39228
7977
```

This process dropped only 772 overall values, and still left approximately 20% of the 1 values in the target category. This should be more than sufficient for our modeling purposes.

Scaling Features

Machine learning models, especially the classification models I will use on this data, are much more effective on standardized data. To transform this dataset into standardized data, I will use the "minmax" scaler from the scikit-learn library.

```
In [28]: # Importing sci-kit minmax scaler
         from sklearn import preprocessing
```

```
In [29]: # Saving column names for future dataframe
         colnames = train.columns
         # Creating minmax scaler instance
         mm_scaler = preprocessing.MinMaxScaler()
         # Transforming data into scaled array
         df_mm = mm_scaler.fit_transform(train)
         # Creating new dataframe with scaled data
         train_clean = pd.DataFrame(df_mm, columns=colnames)
```

```
In [30]: train_clean.describe()
```

Out[30]:

	x0	x1	x2	x3	x4	x5
count	39228.000000	39228.000000	39228.000000	39228.000000	39228.000000	39228.000000
mean	0.477927	0.479577	0.520466	0.514035	0.471174	0.492666
std	0.117946	0.119690	0.120856	0.120769	0.124693	0.123240
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.398198	0.398858	0.438630	0.432242	0.387015	0.409138
50%	0.478068	0.478641	0.521456	0.514728	0.471413	0.493099
75%	0.556340	0.558898	0.601996	0.595674	0.555341	0.575851
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

Now all values have been scaled between the ranges of 0 and 1.

Feature Selection

Now I will narrow down the list of features to help improve the efficiency of the machine learning models. This will involve:

- Checking for correlation with the target column.
- Eliminating features that seem to be unrelated to the target.
- Checking for colinearity to make sure no information is leaking (surviving as a proxy for the target).
- Creating a final dataframe for model testing.

```
In [31]: corr_values = abs(train_clean.corr()['y'])
```

```
In [32]: corr_values.sort_values(ascending=False, inplace=True)  
corr_values
```

```
Out[32]: y      1.000000
         x75     0.204860
         x37     0.198846
         x97     0.187550
         x58     0.184349
         x41     0.176463
         x70     0.106727
         x1      0.104049
         x99     0.099482
         x22     0.098001
         x33     0.096947
         x66     0.096597
         x79     0.096069
         x69     0.095656
         x3      0.093992
         x21     0.092737
         x63     0.092689
         x40     0.092125
         x78     0.091484
         x96     0.091409
         x50     0.090401
         x83     0.090184
         x45     0.089987
         x73     0.089966
         x2      0.089283
         x56     0.088974
         x72     0.088764
         x5      0.087622
         x85     0.086100
         x51     0.086080
         x10     0.082422
         x20     0.080256
         x35_thu 0.073902
         x0      0.067035
         x35_wed 0.058920
         x44     0.054281
         x35_tue 0.049927
         x68_oct 0.026025
         x68_apr 0.025676
         x68_feb 0.022781
         x68_nov 0.022572
         x68_jul 0.021731
         x35_mon 0.018683
         x48     0.013672
         x53     0.013388
         x68_mar 0.013169
         x38     0.012941
         x29     0.011586
         x68_aug 0.011443
         x68_may 0.010876
         x9      0.009982
         x8      0.009883
         x74     0.009397
         x68_jun 0.009169
         x68_dec 0.008615
         x42     0.008502
         x68_sep 0.008000
```


x88	0.007870
x12	0.007306
x92	0.006818
x98	0.006504
x17	0.006355
x46	0.006313
x6	0.006223
x14	0.006138
x7	0.006066
x82	0.005897
x35_fri	0.005893
x30	0.005721
x16	0.005714
x34_chrysler	0.005634
x13	0.005620
x23	0.005538
x18	0.005336
x65	0.005181
x95	0.005096
x52	0.004992
x34_mercedes	0.004813
x34_chevrolet	0.004679
x54	0.004563
x67	0.004432
x62	0.004286
x80	0.004072
x36	0.004032
x31	0.004019
x25	0.003959
x43	0.003343
x60	0.003246
x94	0.002805
x34_toyota	0.002633
x86	0.002397
x57	0.002364
x34_honda	0.002300
x39	0.002262
x89	0.002231
x68_jan	0.002027
x64	0.002010
x47	0.001993
x76	0.001987
x34_volkswagen	0.001826
x27	0.001739
x34_tesla	0.001663
x19	0.001620
x49	0.001581
x55	0.001411
x81	0.001331
x91	0.001320
x90	0.001218
x4	0.001006
x34_ford	0.000929
x24	0.000896
x32	0.000895
x71	0.000841
x15	0.000820

```
x59          0.000712
x34_bmw      0.000710
x11          0.000648
x87          0.000585
x61          0.000513
x34_nissan   0.000336
x26          0.000223
x28          0.000148
x84          0.000028
x77          0.000020
Name: y, dtype: float64
```

It appears as though the values are not strongly correlated with the target variable. (Note: this correlation is being determined with a binary classifier, which is not ideal. Pearson correlation should be used mostly on continuous variables. However, the relative correlation between each feature is what we are interested in, so these results can still be useful for feature selection).


In order to make the model more efficient, I will eliminate any values whose [absolute] correlation is lower than .01

```
In [33]: # Create list of column names with correlations below .01
low_corr_names = list(corr_values[corr_values <= .01].index)
```

```
In [34]: # Dropping list of low correlation features from dataset
train_clean = train_clean.drop(columns=low_corr_names)
```

Now I will plot a correlation heatmap to check for colinearity among any of the remaining features

```
In [35]: # Creating correlation matrix
train_corr = abs(train_clean.corr())
# Mapping correlation matrix
f, ax = plt.subplots(figsize=(20, 20))
ax = sns.heatmap(train_corr)
```




It seems as though the only highly correlated features are 'x35_wed' and 'x35_thu'. These are the dummy variables we created above, so they are not simply representing the same data. However, from the value counts we calculated earlier, the wednesday and thursday categories combined account for over 96% of the data. If these two variables are indeed collinear, then they are representing a larger category that is highly skewed in the wed/thu bucket.

Because of this, we should eliminate the 'x35' feature entirely so it does not throw off the results of the model.

```
In [36]: # Creating list of x35 categories to drop
dropcols = ['x35_mon', 'x35_tue', 'x35_wed', 'x35_thu']
train_clean = train_clean.drop(columns=dropcols)
```

```
In [37]: # Recreating correlation matrix
train_corr = abs(train_clean.corr())
# Mapping correlation matrix
f, ax = plt.subplots(figsize=(20, 20))
ax = sns.heatmap(train_corr)
```



It now appears that collinearity has been eliminated and we have a reduced set of features.

Recursive Feature Elimination

To help improve the model even further, I will use the scikit learn library's recursive feature elimination function to help narrow down the number of significant features for training. Since we are attempting to fit a binary classification model, I will implement this feature elimination using logistic regression with cross-fold validation.

```
In [379]: # Importing feature elimination
from sklearn.feature_selection import RFECV
from sklearn.linear_model import LogisticRegression
```

```
In [380]: # Creating feature and target sets
all_X = train_clean.drop(columns='y')
all_Y = train_clean['y']
```

```
In [383]: # Implementing a logistic regression model
lr = LogisticRegression(max_iter=1000)
# Creating a selector for feature elimination
selector = RFECV(lr, cv=10)
# Fitting the selector and model to training data
selector.fit(all_X, all_Y)
```

```
Out[383]: RFECV(cv=10,
               estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
               fit_intercept=True, intercept_scaling=1,
               l1_ratio=None, max_iter=1000,
               multi_class='auto', n_jobs=None,
               penalty='l2', random_state=None,
               solver='lbfgs', tol=0.0001, verbose=0,
               warm_start=False),
               min_features_to_select=1, n_jobs=None, scoring=None, step=1, verbose=0)
```

```
In [384]: # Checking for optimized columns
Optimized_columns=all_X.columns[selector.support_]
```

```
In [387]: print(Optimized_columns)
          print(len(Optimized_columns))

Index(['x0', 'x1', 'x2', 'x3', 'x5', 'x10', 'x20', 'x21', 'x22', 'x29',
       'x33',
       'x37', 'x38', 'x40', 'x44', 'x48', 'x50', 'x51', 'x53', 'x56',
       'x58',
       'x63', 'x66', 'x69', 'x70', 'x72', 'x73', 'x75', 'x78', 'x79',
       'x83',
       'x85', 'x96', 'x97', 'x99', 'x41', 'x45', 'x68_aug', 'x68_feb',
       'x68_mar', 'x68_nov', 'x68_oct'],
      dtype='object')
42
```

```
In [391]: # finding columns not selected
          not_selected=[]
          total_cols=list(all_X.columns)
          for col in total_cols:
              if col in Optimized_columns:
                  pass
              else:
                  not_selected.append(col)
```

```
In [392]: not_selected
```

```
Out[392]: ['x68_apr', 'x68_jul', 'x68_may']
```

After running this optimization, it seems as though the only features which were eliminated were the april, may and june features. We can eliminate these features from the overall dataset and continue with our machine learning process.

```
In [393]: train_clean = train_clean.drop(columns=not_selected)
```

Model Creation

Error Metric

The error metric I will use for testing the models is provided in the assessment description. Models will be evaluated based on AUC score, so I will use this metric to evaluate the success of each model.

```
In [404]: # Importing auc metric
          from sklearn.metrics import roc_auc_score

          # Import Kfold
          from sklearn.model_selection import KFold
```

Logistic Regression

To begin the Machine Learning process, I will implement a basic logistic regression model, using the features that remain in the 'train_clean' dataset. Logistic regression is a great place to start for a binary classifier, and will give us a baseline against which to test any other models for improvement.

First, I will define a function for the logistic regression model with K-fold cross validation.

```
In [418]: # Def a function for training/testing
def train_and_test(df, k=0):
    # Splitting dataframe from target column
    features = df.columns.drop('y')
    lr=LogisticRegression(max_iter=1000)

    # Building K-folds
    kf = KFold(n_splits=k, shuffle=True)
    auc_values = []
    for train_index, test_index, in kf.split(df):
        # Creating train/test set for fold
        train = df.iloc[train_index]
        test = df.iloc[test_index]
        # Fitting and predicting
        lr.fit(train[features], train['y'])
        predictions = lr.predict(test[features])
        # Calculate AUC
        auc = roc_auc_score(test['y'], predictions)
        auc_values.append(auc)
    # Averaging auc values
    avg_auc = np.mean(auc_values)
    var_auc = np.var(auc_values)
    print(avg_auc, var_auc)
    return avg_auc, var_auc
```

Next, I will train the model with a variety of K-fold values to see which provides the best accuracy.

```
In [415]: # Creating a list of k-fold values
          splits = [2, 4, 6, 8, 10]

          # Creating a dictionary of values for auc and var
          auc_dict = {}

          # Training a model with training data
          for i in splits:
              k_auc, k_var = train_and_test(train, i)
              auc_dict[i] = k_auc, k_var
```



```
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.  
Please also refer to the documentation for alternative solver options:  
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression  
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)  
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.  
Please also refer to the documentation for alternative solver options:  
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression  
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)  
0.7864077223358682 1.445338211400935e-07
```

```
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

```
0.7863791958247519 2.2179188732573277e-05
```

```
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
```

```
https://scikit-learn.org/stable/modules/linear_model.html#logistic-  
regression  
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)  
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/line  
ar_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge  
(status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown
in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.  
Please also refer to the documentation for alternative solver options:  
https://scikit-learn.org/stable/modules/linear_model.html#logistic-  
regression  
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)  
0.784796173973918 5.206692782180382e-05
```

```
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.  
Please also refer to the documentation for alternative solver options:  
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression  
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)  
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.  
Please also refer to the documentation for alternative solver options:  
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression  
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)  
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.  
Please also refer to the documentation for alternative solver options:  
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression  
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)  
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.  
Please also refer to the documentation for alternative solver options:  
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression  
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)  
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.  
Please also refer to the documentation for alternative solver options:
```

```
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/line
ar_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/line
ar_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/line
ar_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

0.7863007769084782 0.00011849641462344976

```
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
```

```

https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

```

https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

```

https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

```

https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

```

https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

```

```
0.7853201730657504 0.00011457379001719168
```



```
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.  
Please also refer to the documentation for alternative solver options:  
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression  
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

```
In [417]: auc_dict
```

```
Out[417]: {2: (0.7864077223358682, 1.445338211400935e-07),  
4: (0.7863791958247519, 2.2179188732573277e-05),  
6: (0.784796173973918, 5.206692782180382e-05),  
8: (0.7863007769084782, 0.00011849641462344976),  
10: (0.7853201730657504, 0.00011457379001719168)}
```

It appears as though the first iteration, with only 2 folds for cross validation, provided the highest AUC score and lowest variance in AUC scores.

Balanced Class Adjustment

To establish a true baseline, however, I need to run this test again using a 'balanced' class model. Since the 1 y-values are disproportionate to the 0 values, this will help balance out the classes and hopefully improve the accuracy of our model.

```
In [419]: # Redefining train/test function for balanced class
def train_and_test(df, k=0):
    # Splitting dataframe from target column
    features = df.columns.drop('y')
    # Including balanced class
    lr=LogisticRegression(max_iter=1000, class_weight='balanced')

    # Building K-folds
    kf = KFold(n_splits=k, shuffle=True)
    auc_values = []
    for train_index, test_index, in kf.split(df):
        # Creating train/test set for fold
        train = df.iloc[train_index]
        test = df.iloc[test_index]
        # Fitting and predicting
        lr.fit(train[features], train['y'])
        predictions = lr.predict(test[features])
        # Calculate AUC
        auc = roc_auc_score(test['y'], predictions)
        auc_values.append(auc)
    # Averaging auc values
    avg_auc = np.mean(auc_values)
    var_auc = np.var(auc_values)
    print(avg_auc, var_auc)
    return avg_auc, var_auc
```

```
In [420]: # Re-running model with balanced classes
# Creating a list of k-fold values
splits = [2, 4, 6, 8, 10]

# Creating a dictionary of values for auc and var
auc_dict = {}

# Training a model with training data
for i in splits:
    k_auc, k_var = train_and_test(train, i)
    auc_dict[i] = k_auc, k_var
```

```
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.  
Please also refer to the documentation for alternative solver options:  
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression  
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)  
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.  
Please also refer to the documentation for alternative solver options:  
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression  
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)  
0.829410606406946 2.219268039043902e-11
```

```
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

```
0.8284743020457778 1.9726083728891866e-05
```

```
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
```

```
https://scikit-learn.org/stable/modules/linear_model.html#logistic-  
regression  
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)  
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/line  
ar_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge  
(status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown
in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.  
Please also refer to the documentation for alternative solver options:  
https://scikit-learn.org/stable/modules/linear_model.html#logistic-  
regression  
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)  
0.8295695189795288 1.3648832633338154e-05
```

```
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
```



```
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/line
ar_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/line
ar_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/line
ar_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

0.8293396643716002 3.137504571509923e-05

```
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
```

```
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/line
ar_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/line
ar_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/line
ar_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/line
ar_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

0.8288388692373765 5.0362273750487695e-05

```
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:939: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html.
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

```
In [421]: auc_dict
```

```
Out[421]: {2: (0.829410606406946, 2.219268039043902e-11),
4: (0.8284743020457778, 1.9726083728891866e-05),
6: (0.8295695189795288, 1.3648832633338154e-05),
8: (0.8293396643716002, 3.137504571509923e-05),
10: (0.8288388692373765, 5.0362273750487695e-05)}
```

It now appears that our best model used 6 K-folds, returning the highest AUC score and still maintaining a low variance. This will be our threshold for improving model accuracy.

Next I will try a K Nearest Neighbors model.

K Nearest Neighbors & Hyperparameter Tuning

K Nearest Neighbors is an effective tool for binary classification models. Here I will implement a basic model and include some hyperparameter tuning through a grid search. This will create several models and evaluate the best hyperparameters to use.

```
In [423]: # Importing K Nearest Neighbors model
from sklearn.neighbors import KNeighborsClassifier
# Importing Grid Search
from sklearn.model_selection import GridSearchCV
```

```
In [424]: # Creating features and target dataframes
all_X = train_clean.drop(columns='y')
all_Y = train_clean['y']
```

```
In [427]: # Dictionary of hyperparameters to test
hyperparameters = {
    "n_neighbors": range(1,20,2),
    "weights": ["distance", "uniform"],
    "algorithm": ['brute'],
    "p": [1,2]
}
# model selection
knn = KNeighborsClassifier()
# grid search cv (param=dictionary we created, cv=folds)
grid = GridSearchCV(knn,param_grid=hyperparameters,cv=10)

grid.fit(all_X, all_Y)

# returning best hyperparameters and score
best_params = grid.best_params_
best_score = grid.best_score_

# training using the best model
best_knn = grid.best_estimator_
```

```
In [429]: best_params
```

```
Out[429]: {'algorithm': 'brute', 'n_neighbors': 5, 'p': 2, 'weights': 'distance'}
```

```
In [430]: # Redefining train/test function for optimized knn model
def train_and_test_knn(df, k=0):
    # Splitting dataframe from target column
    features = df.columns.drop('y')
    # Including balanced class
    knn = KNeighborsClassifier(n_neighbors=5, weights='distance', algorithm='brute', p=2)

    # Building K-folds
    kf = KFold(n_splits=k, shuffle=True)
    auc_values = []
    for train_index, test_index, in kf.split(df):
        # Creating train/test set for fold
        train = df.iloc[train_index]
        test = df.iloc[test_index]
        # Fitting and predicting
        knn.fit(train[features], train['y'])
        predictions = knn.predict(test[features])
        # Calculate AUC
        auc = roc_auc_score(test['y'], predictions)
        auc_values.append(auc)
    # Averaging auc values
    avg_auc = np.mean(auc_values)
    var_auc = np.var(auc_values)
    print(avg_auc, var_auc)
    return avg_auc, var_auc
```

```
In [431]: # Running model for optimized knn  
train_and_test(train_clean, k=10)
```

```
0.814493114647189 5.168933046039276e-05
```

```
Out[431]: (0.814493114647189, 5.168933046039276e-05)
```

KNN Results

After optimizing a K Nearest Neighbor model, the AOC score returned as .8145. This is close to the score of our best Logistic Regression model, but not quite as good.

Random Forest Classifier

The final basic model I will choose is a Random Forest Classifier. This algorithm is well suited to binary classification, and can eliminate the bias of a single decision tree.

I will use the same Grid Search optimization method to try and find the best hyperparameters for a Random Forest model.

```
In [61]: from sklearn.ensemble import RandomForestClassifier
```

```
In [437]: # Dictionary of hyperparameters to test
hyperparameters = {
    "n_estimators": range(100,500,100),
    "max_depth": [10],
    'class_weight': ['balanced'],
    'n_jobs': [-1],
    'verbose': [1]
}
# model selection
rf = RandomForestClassifier()
# grid search cv (param=dictionary we created, cv=folds)
grid = GridSearchCV(rf,param_grid=hyperparameters,cv=10)

grid.fit(all_X, all_Y)

# returning best hyperparameters and score
best_params = grid.best_params_
best_score = grid.best_score_

# training using the best model
best_rf = grid.best_estimator_
```

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks          | elapsed:      5.6s
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:      9.3s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks          | elapsed:      0.0s
[Parallel(n_jobs=4)]: Done 100 out of 100 | elapsed:      0.1s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks          | elapsed:      2.8s
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:      6.5s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks          | elapsed:      0.0s
[Parallel(n_jobs=4)]: Done 100 out of 100 | elapsed:      0.0s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks          | elapsed:      2.9s
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:      7.0s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks          | elapsed:      0.0s
[Parallel(n_jobs=4)]: Done 100 out of 100 | elapsed:      0.0s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks          | elapsed:      2.8s
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:      6.4s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks          | elapsed:      0.0s
[Parallel(n_jobs=4)]: Done 100 out of 100 | elapsed:      0.0s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks          | elapsed:      3.2s
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:      6.8s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks          | elapsed:      0.0s
[Parallel(n_jobs=4)]: Done 100 out of 100 | elapsed:      0.0s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks          | elapsed:      3.0s
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:      7.1s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks          | elapsed:      0.0s
[Parallel(n_jobs=4)]: Done 100 out of 100 | elapsed:      0.0s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks          | elapsed:      3.6s
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:      7.3s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks          | elapsed:      0.0s
[Parallel(n_jobs=4)]: Done 100 out of 100 | elapsed:      0.0s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
```



```

ers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    3.4s
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:    7.0s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 100 out of 100 | elapsed:    0.0s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    2.8s
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:    6.4s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 100 out of 100 | elapsed:    0.0s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    2.8s
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:    6.4s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 100 out of 100 | elapsed:    0.0s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    2.9s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   12.6s
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed:   13.1s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 200 out of 200 | elapsed:    0.1s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    2.8s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   12.2s
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed:   12.7s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 200 out of 200 | elapsed:    0.1s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    2.8s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   12.1s
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed:   12.6s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 200 out of 200 | elapsed:    0.1s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    2.7s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   12.3s

```

```
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed: 12.9s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks | elapsed: 0.0s
[Parallel(n_jobs=4)]: Done 192 tasks | elapsed: 0.1s
[Parallel(n_jobs=4)]: Done 200 out of 200 | elapsed: 0.1s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks | elapsed: 2.8s
[Parallel(n_jobs=-1)]: Done 192 tasks | elapsed: 16.0s
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed: 16.7s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks | elapsed: 0.0s
[Parallel(n_jobs=4)]: Done 192 tasks | elapsed: 0.1s
[Parallel(n_jobs=4)]: Done 200 out of 200 | elapsed: 0.1s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks | elapsed: 3.0s
[Parallel(n_jobs=-1)]: Done 192 tasks | elapsed: 12.8s
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed: 13.2s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks | elapsed: 0.0s
[Parallel(n_jobs=4)]: Done 192 tasks | elapsed: 0.1s
[Parallel(n_jobs=4)]: Done 200 out of 200 | elapsed: 0.1s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks | elapsed: 2.8s
[Parallel(n_jobs=-1)]: Done 192 tasks | elapsed: 12.2s
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed: 12.7s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks | elapsed: 0.0s
[Parallel(n_jobs=4)]: Done 192 tasks | elapsed: 0.1s
[Parallel(n_jobs=4)]: Done 200 out of 200 | elapsed: 0.1s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks | elapsed: 3.6s
[Parallel(n_jobs=-1)]: Done 192 tasks | elapsed: 16.7s
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed: 17.2s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks | elapsed: 0.0s
[Parallel(n_jobs=4)]: Done 192 tasks | elapsed: 0.1s
[Parallel(n_jobs=4)]: Done 200 out of 200 | elapsed: 0.1s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks | elapsed: 3.4s
[Parallel(n_jobs=-1)]: Done 192 tasks | elapsed: 13.7s
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed: 14.3s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks | elapsed: 0.0s
[Parallel(n_jobs=4)]: Done 192 tasks | elapsed: 0.1s
[Parallel(n_jobs=4)]: Done 200 out of 200 | elapsed: 0.1s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
```

ers.

```
[Parallel(n_jobs=-1)]: Done 42 tasks          | elapsed:    4.4s
[Parallel(n_jobs=-1)]: Done 192 tasks         | elapsed:   14.7s
[Parallel(n_jobs=-1)]: Done 200 out of 200    | elapsed:   15.2s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks          | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks         | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 200 out of 200    | elapsed:    0.1s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks          | elapsed:    2.9s
[Parallel(n_jobs=-1)]: Done 192 tasks         | elapsed:   13.5s
[Parallel(n_jobs=-1)]: Done 300 out of 300    | elapsed:   20.5s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks          | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks         | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300    | elapsed:    0.1s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks          | elapsed:    3.1s
[Parallel(n_jobs=-1)]: Done 192 tasks         | elapsed:   15.1s
[Parallel(n_jobs=-1)]: Done 300 out of 300    | elapsed:   23.7s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks          | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks         | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300    | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks          | elapsed:    4.0s
[Parallel(n_jobs=-1)]: Done 192 tasks         | elapsed:   15.8s
[Parallel(n_jobs=-1)]: Done 300 out of 300    | elapsed:   24.2s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks          | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks         | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300    | elapsed:    0.1s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks          | elapsed:    2.9s
[Parallel(n_jobs=-1)]: Done 192 tasks         | elapsed:   15.6s
[Parallel(n_jobs=-1)]: Done 300 out of 300    | elapsed:   29.8s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks          | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks         | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300    | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks          | elapsed:    3.6s
[Parallel(n_jobs=-1)]: Done 192 tasks         | elapsed:   17.9s
[Parallel(n_jobs=-1)]: Done 300 out of 300    | elapsed:   26.4s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks          | elapsed:    0.0s
```

```
[Parallel(n_jobs=4)]: Done 192 tasks      | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300 | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done  42 tasks      | elapsed:    3.9s
[Parallel(n_jobs=-1)]: Done 192 tasks      | elapsed:   14.8s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed:   22.8s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done  42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks      | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300 | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done  42 tasks      | elapsed:    3.1s
[Parallel(n_jobs=-1)]: Done 192 tasks      | elapsed:   14.6s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed:   26.5s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done  42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks      | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300 | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done  42 tasks      | elapsed:    4.9s
[Parallel(n_jobs=-1)]: Done 192 tasks      | elapsed:   16.6s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed:   24.6s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done  42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks      | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300 | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done  42 tasks      | elapsed:    3.5s
[Parallel(n_jobs=-1)]: Done 192 tasks      | elapsed:   14.7s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed:   22.5s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done  42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks      | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300 | elapsed:    0.1s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done  42 tasks      | elapsed:    3.0s
[Parallel(n_jobs=-1)]: Done 192 tasks      | elapsed:   13.6s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed:   21.3s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done  42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks      | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300 | elapsed:    0.1s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done  42 tasks      | elapsed:    3.2s
[Parallel(n_jobs=-1)]: Done 192 tasks      | elapsed:   15.4s
[Parallel(n_jobs=-1)]: Done 400 out of 400 | elapsed:   30.7s finished
```

```
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 400 out of 400 | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    3.1s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   14.3s
[Parallel(n_jobs=-1)]: Done 400 out of 400 | elapsed:   29.3s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 400 out of 400 | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    3.1s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   13.9s
[Parallel(n_jobs=-1)]: Done 400 out of 400 | elapsed:   28.3s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 400 out of 400 | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    3.0s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   13.5s
[Parallel(n_jobs=-1)]: Done 400 out of 400 | elapsed:   28.3s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 400 out of 400 | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    3.2s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   14.1s
[Parallel(n_jobs=-1)]: Done 400 out of 400 | elapsed:   29.5s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 400 out of 400 | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    3.3s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   14.6s
[Parallel(n_jobs=-1)]: Done 400 out of 400 | elapsed:   30.1s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 400 out of 400 | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
```

```

[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    3.3s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   13.9s
[Parallel(n_jobs=-1)]: Done 400 out of 400 | elapsed:   29.2s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 400 out of 400 | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    3.1s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   13.8s
[Parallel(n_jobs=-1)]: Done 400 out of 400 | elapsed:   28.9s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 400 out of 400 | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    3.0s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   14.2s
[Parallel(n_jobs=-1)]: Done 400 out of 400 | elapsed:   29.7s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 400 out of 400 | elapsed:    0.3s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    3.4s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   15.8s
[Parallel(n_jobs=-1)]: Done 400 out of 400 | elapsed:   31.9s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 400 out of 400 | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    3.6s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   15.6s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed:   23.4s finished

```

In [438]: best_params

```

Out[438]: {'class_weight': 'balanced',
           'max_depth': 10,
           'n_estimators': 300,
           'n_jobs': -1,
           'verbose': 1}

```

```
In [441]: # Redefining train/test function for optimized random forest model
def train_and_test_rf(df, k=0):
    # Splitting dataframe from target column
    features = df.columns.drop('y')
    # Including balanced class
    rf = RandomForestClassifier(class_weight='balanced', max_depth=10, n_estimators=300, n_jobs=-1, verbose=1)

    # Building K-folds
    kf = KFold(n_splits=k, shuffle=True)
    auc_values = []
    for train_index, test_index, in kf.split(df):
        # Creating train/test set for fold
        train = df.iloc[train_index]
        test = df.iloc[test_index]
        # Fitting and predicting
        rf.fit(train[features], train['y'])
        predictions = rf.predict(test[features])
        # Calculate AUC
        auc = roc_auc_score(test['y'], predictions)
        auc_values.append(auc)
    # Averaging auc values
    avg_auc = np.mean(auc_values)
    var_auc = np.var(auc_values)
    print(avg_auc, var_auc)
    return avg_auc, var_auc
```

```
In [442]: # Running model for optimized rf  
train_and_test_rf(train_clean, k=10)
```



```

[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    2.5s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   11.2s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed:   17.5s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300 | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    2.6s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   11.3s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed:   17.6s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300 | elapsed:    0.1s finished
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    2.6s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   11.2s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed:   17.6s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300 | elapsed:    0.1s finished
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    2.5s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   11.2s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed:   17.4s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300 | elapsed:    0.1s finished
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    2.6s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   11.3s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed:   17.5s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300 | elapsed:    0.1s finished
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    2.6s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   11.6s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed:   17.8s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.

```

```

[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300 | elapsed:    0.1s finished
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    2.6s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   13.2s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed:   20.8s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300 | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    3.9s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   17.5s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed:   27.5s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300 | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    2.5s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   12.9s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed:   19.9s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s
[Parallel(n_jobs=4)]: Done 300 out of 300 | elapsed:    0.2s finished
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:    2.7s
[Parallel(n_jobs=-1)]: Done 192 tasks     | elapsed:   12.0s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed:   19.2s finished
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent
workers.
[Parallel(n_jobs=4)]: Done 42 tasks      | elapsed:    0.0s
[Parallel(n_jobs=4)]: Done 192 tasks     | elapsed:    0.1s

0.8653249605810925 0.00010615893586465878

[Parallel(n_jobs=4)]: Done 300 out of 300 | elapsed:    0.2s finished

```

Out[442]: (0.8653249605810925, 0.00010615893586465878)

Random Forest Results

This algorithm produced an AOC score of 0.865, significantly higher than previous models. The variance has increased, but not enough to rule out this model as the most effective so far.

Neural Network Application

Last. I will run two Neural Network models to see if the accuracy improves.

```
In [67]: # Importing Neural Network Classifier
from sklearn.neural_network import MLPClassifier
```

```
In [453]: # Redefining neural neetwork train and test function
def train_and_test_nn(df, k=0, n=0):
    # Splitting dataframe from target column
    features = df.columns.drop('y')
    # Including balanced class
    mlp = MLPClassifier(hidden_layer_sizes=(n,), max_iter=500)

    # Building K-folds
    kf = KFold(n_splits=k, shuffle=True)
    auc_values = []
    for train_index, test_index, in kf.split(df):
        # Creating train/test set for fold
        train = df.iloc[train_index]
        test = df.iloc[test_index]
        # Fitting and predicting
        mlp.fit(train[features], train['y'])
        predictions = mlp.predict(test[features])
        # Calculate AUC
        auc = roc_auc_score(test['y'], predictions)
        auc_values.append(auc)
    # Averaging auc values
    avg_auc = np.mean(auc_values)
    var_auc = np.var(auc_values)
    print(avg_auc, var_auc)
    return avg_auc, var_auc
```

```
In [454]: # Running model for different neuron levels in one hidden layer
neurons = [8, 16, 32]
accuracies = {}
for i in neurons:
    aoc_score, aoc_var = train_and_test_nn(train_clean, k=4, n=i)
    accuracies[i] = aoc_score, aoc_var
```

```
0.8333999923376425 0.0006858904796699593
0.8710044161766354 0.0003664845593469178
0.908420581924953 0.0005757786280987429
```

This increased the accuracy level significantly, but also increased the variance of the accuracy levels in the K-fold testing. This could mean the model is over-fitting, but the variance level is still very small compared to the overall accuracy level.

Finally, I will run a Neural Network model with two hidden layers, with 16 nodes each.

```
In [455]: # Redefining neural network for 2 hidden layers
def train_and_test_nn2(df, k=0, n=0):
    # Splitting dataframe from target column
    features = df.columns.drop('y')
    # Including balanced class
    mlp = MLPClassifier(hidden_layer_sizes=(n,n), max_iter=500)

    # Building K-folds
    kf = KFold(n_splits=k, shuffle=True)
    auc_values = []
    for train_index, test_index, in kf.split(df):
        # Creating train/test set for fold
        train = df.iloc[train_index]
        test = df.iloc[test_index]
        # Fitting and predicting
        mlp.fit(train[features], train['y'])
        predictions = mlp.predict(test[features])
        # Calculate AUC
        auc = roc_auc_score(test['y'], predictions)
        auc_values.append(auc)
    # Averaging auc values
    avg_auc = np.mean(auc_values)
    var_auc = np.var(auc_values)
    print(avg_auc, var_auc)
    return avg_auc, var_auc
```

```
In [456]: # Running model for different neuron levels in one hidden layer
neurons = [16]
accuracies = {}
for i in neurons:
    aoc_score, aoc_var = train_and_test_nn2(train_clean, k=4, n=i)
    accuracies[i] = aoc_score, aoc_var
```

0.90822683719581 1.79273849200885e-05

```
In [457]: # Running model for different neuron levels in two hidden layers
neurons = [32]
accuracies = {}
for i in neurons:
    aoc_score, aoc_var = train_and_test_nn2(train_clean, k=4, n=i)
    accuracies[i] = aoc_score, aoc_var
```

0.9302986826327426 1.603054636250812e-05

This model looks to be the best overall, with the highest AOC and one of the lowest variances of the models we have generated.

Generating Predictions

Based on this testing, it seems the two best fitting models are:

- Neural Network Classifier (AOC: 0.9303) **Hyper-parameters: 2 hidden layers, 32 nodes each** K-fold validation: 4 folds
- Random Forest Classifier (AOC: 0.8653) **Hyper-parameters: class_weight: 'balanced', max_depth: 10, n_estimators: 300, n_jobs: -1** K-fold validation: 10 folds

Now I will use these models to generate predictions for the test data and save those predictions in the appropriate format as per the assessment instructions.

Test Data Cleaning

First, I need to clean the test data with the same process used above.

```
In [48]: # Reimporting test data to clean dataframe
test = pd.read_csv('testdata.csv')
```

```
In [49]: # Identifying text columns
text_vals = test.select_dtypes(include=['object'])
text_vals.head()
# Cleaning the values of x41 and x45, storing as float
text_vals['x41'] = text_vals['x41'].str.replace('$', '').astype('float')
text_vals['x45'] = text_vals['x45'].str.replace('%', '').astype('float')
```

```
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy
"""
```

```
/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy
```

```
In [50]: # Creating a mapping dictionary for correct labeling
dict = {'volkswagon': 'volkswagen', 'Toyota': 'toyota',
        'bmw': 'bmw', 'Honda': 'honda', 'tesla': 'tesla',
        'chrystler': 'chrysler', 'nissan': 'nissan',
        'ford': 'ford', 'mercadades': 'mercedes',
        'chevrolet': 'chevrolet'}
# Replacing values using dictionary
text_vals['x34'] = text_vals['x34'].replace(dict)

# Creating mapping dictionary
dict2 = {'July': 'jul', 'Jun': 'jun',
        'Aug': 'aug', 'May': 'may',
        'sept.': 'sep', 'Apr': 'apr',
        'Oct': 'oct', 'Mar': 'mar',
        'Nov': 'nov', 'Feb': 'feb',
        'Dev': 'dec', 'January': 'jan'}
# Replacing values using dictionary
text_vals['x68'] = text_vals['x68'].replace(dict2)
```

/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:8: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

/Users/eddiekirkland/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:18: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

```
In [51]: # Drop x93 from text vals dataframe
text_vals = text_vals.drop(columns='x93')
```

```
In [52]: text_vals.head()
```

Out[52]:

	x34	x35	x41	x45	x68
0	bmw	thursday	107.93	0.00	jun
1	tesla	thursday	-600.43	0.02	may
2	honda	thursday	103.08	-0.00	jun
3	volkswagen	thursday	1518.78	-0.01	sep
4	volkswagen	thursday	-2324.39	-0.00	jun

```
In [53]: # Selecting only text columns
textcols = ['x34', 'x35', 'x68']
# Get dummy prefix names
prefixes = list(text_vals[textcols].columns)
# Getting dummy features for each text column
text_dummies = pd.get_dummies(text_vals, prefix=prefixes)
```

```
In [54]: # Dropping dirty text columns from original dataset
dropcols = ['x34', 'x35', 'x41', 'x45', 'x68', 'x93']
test = test.drop(columns=dropcols)
```

```
In [55]: # Reintroducing our cleaned data into the original dataset
test = pd.concat([test, text_dummies], axis=1)
```

```
In [56]: # Filling null values with mean for feature
test.fillna(value=0, inplace=True)
```

```
In [57]: test.info(verbose=True)
```



```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 10000 entries, 0 to 9999
```

```
Data columns (total 126 columns):
```

```
x0          float64
x1          float64
x2          float64
x3          float64
x4          float64
x5          float64
x6          float64
x7          float64
x8          float64
x9          float64
x10         float64
x11         float64
x12         float64
x13         float64
x14         float64
x15         float64
x16         float64
x17         float64
x18         float64
x19         float64
x20         float64
x21         float64
x22         float64
x23         float64
x24         float64
x25         float64
x26         float64
x27         float64
x28         float64
x29         float64
x30         float64
x31         float64
x32         float64
x33         float64
x36         float64
x37         float64
x38         float64
x39         float64
x40         float64
x42         float64
x43         float64
x44         float64
x46         float64
x47         float64
x48         float64
x49         float64
x50         float64
x51         float64
x52         float64
x53         float64
x54         float64
x55         float64
x56         float64
x57         float64
```

x58	float64
x59	float64
x60	float64
x61	float64
x62	float64
x63	float64
x64	float64
x65	float64
x66	float64
x67	float64
x69	float64
x70	float64
x71	float64
x72	float64
x73	float64
x74	float64
x75	float64
x76	float64
x77	float64
x78	float64
x79	float64
x80	float64
x81	float64
x82	float64
x83	float64
x84	float64
x85	float64
x86	float64
x87	float64
x88	float64
x89	float64
x90	float64
x91	float64
x92	float64
x94	float64
x95	float64
x96	float64
x97	float64
x98	float64
x99	float64
x41	float64
x45	float64
x34_bmw	uint8
x34_chevrolet	uint8
x34_chrysler	uint8
x34_ford	uint8
x34_honda	uint8
x34_mercedes	uint8
x34_nissan	uint8
x34_tesla	uint8
x34_toyota	uint8
x34_volkswagen	uint8
x35_fri	uint8
x35_friday	uint8
x35_monday	uint8
x35_thur	uint8
x35_thursday	uint8

```

x35_tuesday      uint8
x35_wed          uint8
x35_wednesday    uint8
x68_apr          uint8
x68_aug          uint8
x68_dec          uint8
x68_feb          uint8
x68_jan          uint8
x68_jul          uint8
x68_jun          uint8
x68_mar          uint8
x68_may          uint8
x68_nov          uint8
x68_oct          uint8
x68_sep          uint8
dtypes: float64(96), uint8(30)
memory usage: 7.6 MB

```

```

In [58]: # Saving column names for future dataframe
colnames = test.columns
# Creating minmax scaler instance
mm_scaler = preprocessing.MinMaxScaler()
# Transforming data into scaled array
df_mm = mm_scaler.fit_transform(test)
# Creating new dataframe with scaled data
test_clean = pd.DataFrame(df_mm, columns=colnames)

```

```

In [59]: # Matching column names from training dataset
train_cols = train_clean.drop(columns='y')
keepcols = train_cols.columns
test_clean = test_clean[keepcols]

```

Predicting with Random Forest Classifier

```

In [62]: features = train_clean.drop(columns='y')

# Running Model on training data
rf = RandomForestClassifier(class_weight='balanced', max_depth=10, n_estimators=300, n_jobs=-1, verbose=1)
rf.fit(features, train_clean['y'])
predictions = rf.predict_proba(test_clean)

```

```
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 4 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 42 tasks | elapsed: 2.8s
```

```
[Parallel(n_jobs=-1)]: Done 192 tasks | elapsed: 12.0s
```

```
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed: 19.0s finished
```

```
[Parallel(n_jobs=4)]: Using backend ThreadingBackend with 4 concurrent workers.
```

```
[Parallel(n_jobs=4)]: Done 42 tasks | elapsed: 0.0s
```

```
[Parallel(n_jobs=4)]: Done 192 tasks | elapsed: 0.1s
```

```
[Parallel(n_jobs=4)]: Done 300 out of 300 | elapsed: 0.2s finished
```

```
In [63]: # Saving 1 predictions to dataframe  
rf_predictions = pd.DataFrame(predictions[:,1])
```

```
In [65]: # Saving Random Forest predictions to file  
rf_predictions.to_csv('RFPredictions_EKirkland.csv')
```

Predicting with Neural Network Classifier

```
In [69]: features = train_clean.drop(columns='y')  
  
# Running Model on training data  
mlp = MLPClassifier(hidden_layer_sizes=(32,32), max_iter=500)  
mlp.fit(features, train_clean['y'])  
predictions = mlp.predict_proba(test_clean)
```

```
In [70]: # Saving 1 predictions to dataframe  
nn_predictions = pd.DataFrame(predictions[:,1])  
# Saving Neural Network predictions to file  
nn_predictions.to_csv('NNPredictions_EKirkland.csv')
```